

# Vectorization in Fortran and C

Eric J. Kostelich



ARIZONA STATE UNIVERSITY  
SCHOOL OF MATHEMATICAL & STATISTICAL SCIENCES

Sept. 19, 2018

# Outline of the lecture today

- Vectorization is form of **data-parallel computation**
- Vector constructs in MATLAB
- Vector constructs in Fortran

## Review: The vectorization paradigm

- **SIMD**: Single Instruction Multiple Data
- To be data parallel, the results of a computation **must not depend** on the order in which the operands are processed
- **Example**: Vector assignment is data parallel:  
$$y(:) = x(:)$$

# Memory hierarchies for the Intel i7 Core processor

RAM (4,096+ MB)  
latency: ~250 cycles

L3 cache (6 MB)  
latency: 40 cycles

L2 cache  
1/4 MB, 12 cycles

L1 cache  
1/32 MB, 3 cycles

## Example: Cache access considerations, $n \times 2$ arrays

$y(:,1) = \text{param.a} - x(:,1).^2 + \text{param.b} * x(:,2); y(:,2) = x(:,1);$

- The **working set** is  $x(1:n,1)$ ,  $x(1:n,2)$ ,  $y(1:n,1)$ ,  $y(1:n,2)$
- Caches operate in last-in first-out (LIFO) order
- The L2 cache holds 256 KB  $\rightarrow$  64 KB per array slice, or a maximum of  $\sim 8,000$  double-precision numbers each
- If  $n > 8,000$  then the array slices **spill** from the L2 cache

## Cache access considerations for multicore programming

- On a typical chip, all cores share the same L3 cache (6–12 MB)
- Usually this equals about 1.5 MB of L3 cache per core
- You get good performance in multithreading only if the working set fits in this space
- Otherwise, memory accesses go to main memory, where the stalls get longer as more cores request data
- Memory constraints limit the performance improvement that can be expected with multithreading

## Additional comments on data structures

- **Example:** Computational geometry

Common textbook usage

type point

real x, y, z

end type point

type(point) points(n)

Superior data structure:

real points(3,n)

- Rotation by **R**: `points=matmul(R, points)`
- Matrix multiplication has highly optimized implementations

# Vectorization in C/C++

- Can this loop be vectorized?

```
void sub(float *x, float *y, size_t n) {  
    for(size_t j = 0; j < n; j++)  
        y[j] = x[j];  
}
```



# Vectorization in C/C++

- Can this loop be vectorized?

```
void sub(float *x, float *y, size_t n) {  
    for(size_t j = 0; j < n; j++)  
        y[j] = x[j];  
}
```

- In general, **no**
- The order of operations matters if **x** and **y** point to overlapping sections of memory

## Vectorization in C99, 2

- Many other optimizations depend on knowing whether  $x$  and  $y$  overlap in memory
- The C99 keyword `restrict` asserts that the pointers aren't aliased

```
void sub(float * restrict x, float * restrict y, size_t n) {  
    for(size_t j = 0; j < n; j++)  
        y[j] = x[j];  
}
```

- This loop vectorizes—but be sure that  $x$  and  $y$  don't point to overlapping regions of memory!
- C++ does not have `restrict` (an important incompatibility with C)

# Vectorization using the C++ STL

```
void saxpy(double a, vector<double>& x, vector<double>&
y) {
    vector<double>::size_type n = x.size();
    for(vector<double>::size_type j = 0; j < n; j++)
        y[j] += a*x[j];
}
```

- Be sure to pass the arguments by reference—otherwise C++ passes a copy and the routine **has no effect**
- Intel's `icpc` vectorizes
- GNU `g++` does not, even with `-O3 -ftree-vectorize`

## Vectorization using the C++ STL, 2

```
void  
saxpy(double a, valarray<double>& x, valarray<double>& y)  
{  
    y *= a*x;  
}
```

- This code does **not** work with **vector**
- **valarray** defines helper classes to make such expressions efficient without creating temporaries
- Intel's **icpc** vectorizes this construct but GNU **g++** does not

# Vectorization in Fortran

```
subroutine saxpy(a, x, y, n)
integer, intent(in):: n
real, intent(in):: a, x(n)
real, intent(inout):: y(n)
y = y + a*x
```

- **Fortran rule:** subroutine arguments **must not refer** to overlapping sections of memory
- Compilers often cannot diagnose violations of the rule
- But if you violate the rule, then you'll get what you deserve

# MATLAB and Fortran syntax

- Suppose  $\mathbf{A} \in \mathbb{R}^{m \times n}$ ,  $\mathbf{x} \in \mathbb{R}^n$ , and  $\mathbf{y} \in \mathbb{R}^m$
- **MATLAB and Fortran:**  $\mathbf{A}(:,k)$  refers to the  $k$ th column of  $\mathbf{A}$  and  $\mathbf{A}(k,:)$  to the  $k$ th row
- **MATLAB and Fortran:**  $\mathbf{A}(:,k)$  is much more efficient (stride 1 memory access)
- **MATLAB:**  $\mathbf{y} = \mathbf{A} * \mathbf{x}$  yields  $\mathbf{y} = \mathbf{Ax}$
- **Fortran:**  $\mathbf{y} = \text{matmul}(\mathbf{A}, \mathbf{x})$  yields  $\mathbf{y} = \mathbf{Ax}$

## MATLAB and Fortran syntax, 2

- Suppose  $\mathbf{x}, \mathbf{y} \in \mathbb{R}^n$  and  $a \in \mathbb{R}$
- **MATLAB and Fortran:**  $\mathbf{y} = \mathbf{y} + a*\mathbf{x}$  is a SAXPY
- Alternative syntax:  $\mathbf{y}(:) = \mathbf{y}(:) + a*\mathbf{x}(:)$
- Whether you write  $\mathbf{x}$ ,  $\mathbf{x}(:)$ ,  $A$ , or  $A(:, :)$  does not matter in MATLAB or in most Fortran usage

## MATLAB and Fortran syntax, 3

- To add  $\mathbf{x}$  to each column of the  $n \times n$  matrix  $\mathbf{A}$ :

- **MATLAB #1:**

```
for k=1:n
```

```
    A(:,k) = A(:,k) + x;
```

```
end
```

- **MATLAB #2:** Add a  $1 \times n$  tiling of  $\mathbf{x}$  to  $\mathbf{A}$ :

```
A = A + repmat(x,1,n);
```

- Either alternative is faster than

```
for k=1:n; for j=1:n;
```

```
    A(j,k) = A(j,k) + x(j);
```

```
end; end
```



## MATLAB and Fortran syntax, 4

- Although vectorized, this code is inefficient in both MATLAB and Fortran:

```
for k=1:n  
    A(k,:) = A(k,:) + x(k);  
end
```

- If  $n$  is greater than the cache line size, then successive memory accesses must wait for main memory
- The performance loss can be up to a factor of 80 for Intel Core i7 processors

## MATLAB and Fortran syntax, 5

- MATLAB and Fortran syntax is identical for the elementwise operations  $+$  and  $-$
- $*$ ,  $/$ , and  $**$  are elementwise in Fortran:

do j=1,n

c(j) = a(j) \* b(j)

enddo

Equivalent assignment:

c = a \* b

- MATLAB requires  $c = a .* b$ ; (and analogously  $./$  and  $.^$ )
- To solve  $\mathbf{Ax} = \mathbf{b}$  in MATLAB:  $x = A \setminus b$ ; (no Fortran equivalent)
- Remember:  $A*x$  is matrix multiplication in MATLAB

## Comments on `matmul`

- `c = matmul(a,b)` works for matrix-vector and matrix-matrix multiplication
- Remember:  $\mathbf{A}_{m \times n} \mathbf{B}_{n \times k} = \mathbf{C}_{m \times k}$
- The result of `matmul` is undefined if the dimensions of the operands aren't compatible
- Assuming compatible dimensions,  $\mathbf{C} = \mathbf{A}^T \mathbf{B}$  is `C=matmul(transpose(A),B)` in Fortran (and is  $\mathbf{C} = \mathbf{A}' * \mathbf{B}$  in MATLAB)

# Common matrix-vector functions with optimized implementations

- **Dot product:** MATLAB:  $\text{dp} = \text{dot}(\mathbf{x}, \mathbf{y})$ ; Fortran:  $\text{dp} = \text{dot\_product}(\mathbf{x}, \mathbf{y})$
- **Arithmetic mean:** MATLAB and Fortran:  $\text{avg} = \text{sum}(\mathbf{x})/n$
- **Maximum value in MATLAB:** 1-d vector:  $\text{big} = \text{max}(\mathbf{x})$ ; 2-d vector:  $\text{big} = \text{max}(\text{max}(\mathbf{x}))$ , etc
- **Maximum value in Fortran:** All ranks:  $\text{big} = \text{maxval}(\mathbf{x})$
- Fortran:  $\mathbf{z} = \text{max}(\mathbf{x}, \mathbf{y})$  returns  $\text{max}(x_i, y_i)$  elementwise
- Analogously for minimum values

# Vectorized loops in Fortran

- A big advance in compiler technology in the 1970's was the advent of vectorizing Fortran compilers
- Loop constructs like

do j=1,n

$y(j) = y(j) + a * x(j)$

enddo

are converted automatically to the appropriate sequence of vector instructions

- There is no performance difference with  $y = y + a * x$ , except that the latter is more concise

# Requirements for vectorizability in Fortran

- Must be a counted **do** loop
- The result must not depend on the order of operations
- Any **if** statements must be simple (no **else if** clauses)
- No other branches within or out of the loop are permitted

Not vectorizable:

```
do j=1,n
```

```
    a(j) = a(j-1) + b(j)
```

```
enddo
```

Vectorizable:

```
do j=1,n
```

```
    if(x(j) > 0.0) x(j)=log(x(j))
```

```
enddo
```

## Requirements for vectorizability, 2

- The result must not depend on the order of operations
- **Consequence:** The destination must not overlap with the source!

Not vectorizable:

```
do j=1,n  
  x(j)=x(n-j+1)  
enddo
```

Vectorizable, if  $x$  and  $y$  don't overlap:

```
do j=1,n  
  y(j)=x(n-j+1)  
enddo  
y=x
```

# Use array constructs with care

Legal Fortran construct:

$x=x(n:1:-1)$

Equivalent to:

do  $j=1,n$

$y(j)=x(n-j+1)$

enddo

$y=x$

- The array-valued assignment **requires** the compiler to generate a temporary array
- The assignment is **as if** the right-hand side were evaluated at once, **then** assigned to the left-hand side



# Masked assignment in Fortran

- Given `real:: x(n)`, the following two constructs are equivalent:

Vectorizable loop:

```
do j=1,n
  if(x(j) > 0) x(j)=log(x(j))
enddo
```

Masked assignment:

```
where(x > 0) x=log(x)
```

- Masked assignment with alternative:

```
where(x > 0)
  x=log(x)
elsewhere
  x= -huge(x)  most negative possible value
endwhere
```

## Masked assignment in Fortran, 2

- The **where** statement is a block construct:

```
where(x > 0)
  x=log(x)
  y=x*log(x)
endwhere
```

- and it can be nested:

```
where(x > 0)
  y=log(x)
  where(y > 0) z=x*log(y)
endwhere
```

# Masked assignment in MATLAB

- The constructs apply to arrays of any rank
- **Elegant:**  $x(x > 0) = 0$  is equivalent to  $x = \min(x, 0)$
- But nesting is clunky:

Fortran:

where( $x > 0$ )

$y = \log(x)$

    where( $y > 0$ )  $z = x * \log(y)$

endwhere

MATLAB:

$ix = \text{find}(x > 0);$

$y(ix) = \log(x(ix));$

$iy = \text{find}(y(ix) > 0);$

$z(ix(iy)) = \dots$

$x(ix(iy)) * \log(y(ix(iy)))$